

Report on the Probabilistic Language Scheme

Alexey Radul
Massachusetts Institute of Technology
77 Massachusetts Avenue
Cambridge, MA 02139
axch@mit.edu

ABSTRACT

Reasoning with probabilistic models has become a widespread and successful technique in areas ranging from computer vision, to natural language processing, to bioinformatics. Currently, these reasoning systems are either coded from scratch in general-purpose languages or use formalisms such as Bayesian networks that have limited expressive power. In both cases, the resulting systems are difficult to modify, maintain, compose, and interoperate with. This work presents Probabilistic Scheme, an embedding of probabilistic computation into Scheme. This gives programmers an expressive language for implementing modular probabilistic models that integrate naturally with the rest of Scheme.

1. INTRODUCTION

Some of the most challenging tasks faced by computers today involve drawing conclusions from noisy or ambiguous data. These tasks range from deciding whether an email message is spam based on the words it contains [10], to predicting whether a piece of ground is safe to drive on based on camera and laser range-finder readings [14], to discovering patterns of gene expression based on microarray data [11]. Probabilistic modeling has become the technique of choice for tackling many of these tasks (as illustrated by the papers just cited). Probability theory provides a well-understood mathematical framework for combining multiple sources of evidence, but the computational tools available to programmers wishing to avail themselves of it are limited.

Probabilistic Scheme is a library for implementing probability models in Scheme. Probabilistic Scheme deals with models of phenomena in arbitrarily structured but discrete and countable possibility spaces. The central concept is that of a probability distribution. A probability distribution, in this context, is a belief about the value that an expression could have when evaluated in some environment. Suppose, for example, that I were about to roll a mathematically perfect six-sided die. Then you should believe that each of the six faces is equally likely to come up as the final result, and

no other results are possible. This belief is a **probability distribution**, which assigns probability one-sixth to each of the faces of the die. Suppose, then, that you asked me the parity of the result and I informed you that the number rolled on the die were odd. Then, assuming you had no doubts about my perceptiveness or veracity, your belief about the result of the roll would change to one that assigned probability one-third to each of the faces with odd numbers, and zero to the others.

Probabilistic Scheme permits one to represent such beliefs, from the simple to the complex; to compute their transformations; and to extract definite, quantitative information from them. As a preview, the beliefs in the preceding paragraph can be represented as follows

```
(define die-roll-distribution
  (make-discrete-distribution
    '(1 1/6) '(2 1/6) '(3 1/6)
    '(4 1/6) '(5 1/6) '(6 1/6)))

(define odd-die-roll-distribution
  (conditional-distribution
    die-roll-distribution odd?))

(distribution/determine! odd-die-roll-distribution)
(distribution/datum-probability
  odd-die-roll-distribution 1)
;Value: 1/3
```

Conceptually, distributions in Probabilistic Scheme are lists of the possibilities, and the probabilities assigned to them. If some object does not appear in the list, its probability is zero. To permit large (or infinite) distributions which are not needed in their entirety at any one time, this is actually a stream¹ of possibilities. The representation of probability distributions will be further detailed in Section 5 and Section 6.

Probabilistic Scheme offers two “languages” for creating and transforming distributions. There is a language for constructing probability distributions by writing nondeterministic Scheme programs, described in Section 3. There is also a language for explicitly creating and manipulating objects that represent probability distributions, described in

¹A stream is a (possibly infinite) list whose actual construction has been delayed so that it is only ever computed as far as any client requests [1].

Section 4. The explicit language is easier to think about, because there are no weird control structures or strange non-deterministic programs, but the nondeterministic language is better suited to defining complex distributions, because it makes the structure of the distribution much clearer and its expression much more natural.

Probabilistic Scheme also offers a language for querying distributions to extract definite information from them. This turns out to be somewhat nontrivial, and is discussed in Section 5. We discuss our implementation of these languages in Section 6, exemplify Probabilistic Scheme in Section 7, and summarize in Section 8. First, though, some background, in Section 2.

2. BACKGROUND

Bayesian networks have been a mainstay of probabilistic modeling since the late 1980s, e.g. [9]. They are a wonderful basic tool, and well implemented by off-the-shelf inference packages such as BUGS [12], but they are not very good at capturing the structure present in a domain. Recent work has moved in the direction of increased structure, for instance capturing relational structure as in [5] and [3], or first-order logical structure as in [7]. Probabilistic Scheme goes beyond this work, in the sense that programs in a general-purpose programming language are as structured as it is possible to be.

The closest modern relative to Probabilistic Scheme is a stochastic programming language based on OCaml, called IBAL [8]. Probabilistic Scheme differs from IBAL in being an embedding of inference into an existing programming language, rather than a programming language in its own right, permitting Probabilistic Scheme to benefit from all of the extant Scheme constructs. The stochastic distribution sub-language of Probabilistic Scheme owes a great intellectual debt to McCarthy’s `amb` operator (see, e.g. [6], [1]).

3. STOCHASTIC LANGUAGE

Probabilistic Scheme embeds probabilistic computation by allowing Scheme expressions to have uncertain values, and maintaining an implicit probability distribution over what those values might be. The primitives for handling implicit distributions are

- `discrete-select` introduces uncertainty
- `observe!` constrains the implicit distribution
- `stochastic-thunk->distribution` encloses a non-deterministic computation and returns the implicit distribution explicitly.

We now discuss each of these primitives in detail.

`(discrete-select possibility ...)`

Takes any number of literal two-element lists representing object-probability pairs. Returns one of the objects, implicitly distributed according to the distribution specified by the probabilities, which are expected to sum to 1. The evaluation of each object is deferred until it needs to be returned.

As expressions combine, their implicit distributions transform according to the rules of probability theory. For example, we can

```
(define (roll-die)
  (discrete-select (1 1/6) (2 1/6) (3 1/6)
                  (4 1/6) (5 1/6) (6 1/6)))
```

Then every call to the `(roll-die)` function will independently return one of the numbers from 1 through 6, implicitly uniformly distributed. In that case, the expression `(cons (roll-die) (roll-die))` returns one of the 36 cons cells that have one of those numbers in the car slot and one in the cdr slot, also implicitly uniformly distributed. The expression `(+ (roll-die) (roll-die))` will return one of the numbers from 2 through 12, implicitly distributed according to the probability of getting that sum when rolling two fair six sided dice.

These rules of combination allow one to define arbitrarily complex distributions over arbitrarily structured objects. A note of caution: The stochastic language plays best with the functional subset of Scheme. Since `discrete-select` causes the implementation of Probabilistic Scheme to run various chunks of code fewer or more times than once, the results of mixing `discrete-select` with side-effects are unspecified.

`(observe! boolean)`

Modifies the current implicit distribution by conditioning it on the argument being true. Returns an unspecified value.

Consider, for example, the expression

```
(let ((face (roll-die))) ;; Line 1
  (observe! (> face 2)) ;; Line 2
  face)                  ;; Line 3
```

In line 1, the expression `(roll-die)` returns one of the numbers from 1 through 6, implicitly uniformly distributed. `let` then binds it to the name `face`, whose value is then implicitly uniformly distributed over 1 through 6. The expression `(> face 2)` on line 2 has one of the values `#t`, `#f`, implicitly distributed as 2/3 for `#t` and 1/3 for `#f`. `Observe!` modifies this implicit distribution to require `#t`. This modifies the implicit distribution for `face` to be consistent with `(> face 2)` returning `#t`, that is it conditions $p(\text{face})$ on `(> face 2)`. The distribution of return values from this whole `let` form is then $p(\text{face} | (> \text{face } 2))$, in other words uniform over the numbers from 3 through 6.

`(stochastic-thunk->distribution thunk)`

Returns, as an explicit probability distribution, the implicit distribution over the possible return values of the given thunk (nullary procedure).

For the example above, `(stochastic-thunk->distribution roll-die)` would return an explicit distribution object that

represented the distribution that assigns equal mass to the numbers 1 through 6. `stochastic-thunk->distribution` captures and contains the nondeterminism occurring inside its argument thunk and perfectly deterministically returns an object representing a probability distribution.²

One way to think of the semantics of the stochastic language is to imagine a possible implementation by rejection sampling. One could implement rejection sampling for this stochastic language by having `discrete-select` just use a random number generator, having `observe!` raise a distinguished exception if its argument is false, and having `stochastic-thunk->distribution` run its thunk many times, recording every successful return as a sample, and throwing away runs that raised the exception. In the limit of running the thunk an infinite number of times, the samples thus produced would constitute the distribution defined by such a program. The actual implementation systematically searches the space of possibilities instead, but that discussion belongs in Section 6.

4. EXPLICIT DISTRIBUTION LANGUAGE

As well as specifying distributions with nondeterministic thunks given to `stochastic-thunk->distribution`, one can create and operate on explicit probability distributions directly. The primitives³ for handling explicit distributions are

- `make-discrete-distribution` creates an explicit probability distribution from a list of possibilities
- `dependent-product` combines two distributions
- `conditional-distribution` transforms a distribution by conditioning it on a predicate
- `distribution-select` makes an explicit distribution implicit

(`make-discrete-distribution` possibility ...)

Interprets each possibility argument as a two-element list of an object and its probability. Returns the probability distribution that assigns those probabilities to those

²Nondeterministic computations can be nested:

```
(stochastic-thunk->distribution
 (lambda ()
  (discrete-select
   ((stochastic-thunk->distribution
    (lambda ()
      (discrete-select ('heads 1/2)
                        ('tails 1/2)))) 1/5)
   ((stochastic-thunk->distribution
    (lambda ()
      (discrete-select
       (1 1/3) (2 1/3) (3 1/3)))) 4/5))))
```

will return an explicit probability distribution weighted 1/5 to 4/5, whose two data are two different explicit probability distributions.

³Actually, `distribution-select` is enough for everything, because the stochastic language can readily represent these operations. This list is primitive in the sense that these operations are a sufficient base even without reference to the stochastic language.

objects, and zero to all others. Expects the set of possibilities to be normalized, i.e. for the given probabilities to sum to 1.

(`dependent-product` distribution function combiner)

A distribution $p(y|X)$ that depends on the value of some variable X can be represented as a function of X that, when given any particular value x , returns the distribution $p(y|X = x)$. Given a distribution $p(x)$ and such a function `(lambda (x) p(y|X=x))`, `dependent-product` returns the distribution $p(x, y) = p(x)p(y|X = x)$. Instead of trying to represent a distribution over multiple values, `dependent-product` takes a combiner to apply to the values x and y , to return $p(\text{combiner } x \ y)$. An oft-useful combiner is `cons`, though the right summation will happen if the combiner maps multiple pairs to the same combined value.

(`conditional-distribution` distribution predicate)

Given a distribution $p(x)$ and a predicate $A(x)$, returns the distribution over x 'es that satisfy the predicate, which is given by

$$p(x|A(x)) = \begin{cases} p(x)/p(A) & \text{if } A(x) \text{ is true} \\ 0 & \text{if } A(x) \text{ is false} \end{cases}$$

where $p(A)$ is the probability that A is true. Since the x 'es are mutually exclusive and exhaustive, we know that

$$p(A) = \sum_{x:A(x)} p(x).$$

The behavior of the system is unspecified if the predicate $A(x)$ is impossible to satisfy. The present implementation will naturally tend either to enter an infinite loop if the underlying stream is infinite, or fail with a divide-by-zero error if it is finite.

(`distribution-select` distribution)

Returns one of the possible values from the given explicit distribution, implicitly distributed according thereto.

For example, `roll-die`, above, could have been defined as

```
(define (roll-die)
 (distribution-select
  (make-discrete-distribution
   '(1 1/6) '(2 1/6) '(3 1/6)
   '(4 1/6) '(5 1/6) '(6 1/6))))
```

5. QUERYING DISTRIBUTIONS

Probability distributions are occasionally infinite, as for instance a Poisson distribution over the integers, and often technically finite but large enough that we do not wish to compute them fully, as for instance a distribution over the possible parses of some sentence. Consequently, Probabilistic Scheme sports a lazy representation of distributions. On initial creation, a distribution is a completely unforced stream

of possibilities, each of which names a value and some amount of probability that the distribution assigns to it. On user request, this internal stream can be (partially) forced, whereupon the distribution object caches the assignments that came out of it. Therefore, at any one time, a distribution will have some cache of the possibilities that came out of the stream so far, which it can use to answer questions, and an upper bound on the amount of probability remaining in the rest of the stream.

It is convenient to permit the internal stream of a distribution to emit impossibilities as well as possibilities. An impossibility has the meaning that some amount of probability “vanishes,” in which case the distribution object implicitly renormalizes. This happens, for instance, if some `observe!` statement forces some condition that otherwise had some probability of being false.

5.1 Questions

The following queries can be applied to explicit probability distributions without further forcing:

`(distribution? thing)`

Returns `#t` if the given thing is an object explicitly representing a probability distribution, and `#f` otherwise.

`(distribution/determined? distribution)`

Returns whether the given distribution object has already been fully determined, as opposed to having more computation it could do to further refine its internal representation of the distribution it represents.

`(distribution/undetermined-mass distribution)`

Returns the amount of probability mass that remains in the unforced segment of the internal possibility stream in this distribution.

`(distribution/datum-min-probability distribution datum)`

`(distribution/datum-max-probability distribution datum)`

`(distribution/datum-probability distribution datum)`

Return bounds on the probability that the given datum could have in this distribution. The minimum value will be realized if all the remaining undetermined mass goes to other data. The maximum value will be realized if all the remaining undetermined mass goes to this datum. The unqualified function will signal an error if any undetermined mass remains, because then the probability of the datum is as yet unknown.

5.2 Forcing

The following functions cause explicit probability distribution objects to perform more of their computations:

`(distribution/refine! distribution)`

Runs the computation in the given distribution for the smallest detectable increment, which is either until a possibility is discovered or until some undetermined mass is

lost to an impossibility. If the former comes to pass, the `min-probability` of the datum of the discovered possibility increases, and the `max-probability` of every other datum decreases. In the latter case, the `min-probability` of every thus far discovered datum increases and the `max-probability` of every datum decreases, unless there could be only one datum. The `undetermined-mass` decreases unless no data have yet been found, in which case it remains 1.

If the distribution computation has been finished, i.e. no undetermined mass remains, `distribution/refine!` does nothing and returns `#f`. If `distribution/refine!` changed something, it returns `#t`. Higher-level forcing functions can be built by iterating `distribution/refine!` for some desired amount of time or until some desired condition has been met.

`(distribution/determine! distribution)`

Runs the computation in the given distribution all the way to the end. This is useful primarily for testing.

`(distribution->density-stream distribution)`

Returns a stream of the possibilities in the given distribution. The stream is permitted to contain impossibilities and repeated data at the discretion of the underlying implementation. This is an effective way to iterate over all the possibilities of a distribution, without requiring it to compute any beyond those that the client deems interesting. The returned stream starts with values from the distribution’s cache, but will begin to force the distribution’s internal stream when necessary (which forcing will be correctly cached for future access).

6. IMPLEMENTATION

We first discuss our implementation of the explicit probability distribution objects, and then proceed to the implementation of the stochastic language.

6.1 Distributions

As mentioned in Section 5, a probability distribution is fundamentally a stream of possibilities and impossibilities. A possibility assigns some density to some particular datum, and an impossibility asserts that some density “disappears”, and the rest should be renormalized to account for that.⁴ Our distribution objects do not perform the renormalization eagerly, but instead track the total density of the impossibilities, and perform the renormalization on the fly as clients ask for the probabilities of various data.

A distribution is then a record containing four components:

- a lazy stream of the possibilities and impossibilities that form the distribution;
- a hash table mapping the data that we have encountered so far to their respective densities;

⁴We use the word *density* here instead of probability to emphasize the need for normalization.

- a measure of how much density remains in the stream, the *undetermined density*; and
- the total density of the impossibilities we have encountered, the *discarded density*.

The probability of some datum in a given distribution is its normalized density — density divided by the distribution’s *normalization constant*, which is one minus the density that has been discarded due to impossibilities. If a distribution has not been completely determined, i.e. the possibility stream has not been exhausted, the probability of any datum is not completely certain. It can, however, be bounded from above by assuming that all the undetermined density will go to this datum, and from below by assuming it will all go to other data. Consequently, a distribution object can compute the bounds on the probability of a given datum in constant time.

The lazy stream allows us to incrementally refine the distribution. As the stream is forced, the elements are recorded in the distribution; possibilities are stored in the hash table, and densities of impossibilities are added to the discarded density field.

The distribution object described above is effective for answering questions about distributions. To implement distribution transformations, such as **dependent-product** and **conditional-distribution**, we only need the raw representation as streams of possibilities. **Conditional-distribution** is particularly simple, in that all it needs to do is check each datum against the predicate, and replace it with an impossibility of the same density if the predicate rejects it.

6.2 Stochastic Language

A program written with **discrete-select** defines a search tree of possible values that calls to **discrete-select** could return. Probabilistic Scheme systematically searches this tree to produce a stream of possible result values and their probabilities.

This search is actually implemented by **discrete-select** capturing its return continuation using one invocation of **call-with-current-continuation**⁵ and saving it in a schedule of branch points to be explored. **Discrete-select** likewise saves the possible options and their given probabilities in this schedule. Exploring an edge in the search tree consists of asking the schedule to pick a saved branch point and option, and escape into that continuation with that value. The computation then proceeds normally until it reaches another **discrete-select** or until it returns a value to the enclosing **stochastic-thunk->distribution**. Said enclosing **stochastic-thunk->distribution** can then return that value to the client, and suspend the search until another value is requested.

During the progress of this search, Probabilistic Scheme maintains the probability (actually density, since observations in other branches do not eagerly renormalize it) of

⁵This is a wonderful but mindbending Scheme control construct, which this footnote lacks the space to explain. It is defined, for instance, in the Scheme Report [4], and explanations and tutorials abound on the Web.

reaching the current point in the program. **Discrete-select** saves this density along with its continuation, and when returning an option, updates it to be the density for getting to that choice point times the probability of choosing that option once there. This is the bookkeeping necessary to allow the search to yield possibilities and impossibilities with associated densities.

Calls to **observe!** either do nothing if the observed condition happens to be true in the current branch, or abort consideration of the current branch if the observed condition proves false, supplying an impossibility to the enclosing **stochastic-thunk->distribution**.

The implementation of **stochastic-thunk->distribution** sets up the background state necessary to execute a search through such a tree (such as the escape continuation for detecting an impossibility) and lazily launches the process, returning a stream of the possibilities and impossibilities the search will discover.

Distribution-select is just like **discrete-select**, except that it derives its list of options from the given distribution instead of from an explicit list in its arguments.

7. EXAMPLES

For our first example, consider the definitions in Figure 1. The **geometric-select** function will return some integer greater than or equal to **start**, implicitly distributed according to a geometrically receding distribution parametrized by **alpha**. Limitations of time and computer memory aside, this distribution is infinite, but the lazy nature of the implementation permits the call to **stochastic-thunk->distribution** to return a perfectly good distribution object. One could then call **distribution/refine!** on it until satisfaction, and **distribution/min-probability** would tell one that 0 has probability at least 1/4 in this distribution. The upper bound, returned by **distribution/max-probability**, would decrease as one refined the distribution further and further, tending to 1/4 in the limit.

The infinite nature of the distribution does no harm to compositional operations either. For instance, one could require that the integers be odd either explicitly with

```
(conditional-distribution receding-distribution odd?)
```

or implicitly with

```
(let ((number (geometric-select 1/4 0)))
  (observe! (odd? number))
  ...)
```

In this situation, the **min-probability** of 0 would remain 0 forever, because it is ruled out by the predicate **odd?**, and both the **min** and **max** probabilities of 1 would tend, as they should, to 7/16, from below and above, respectively, as one refined the resulting distribution further.

The fun doesn’t end there! Figures 2, 3 and 4 exemplify definitions of distributions over arbitrary, structured objects.

```
(define (geometric-select alpha start)
  (discrete-select
    (start alpha)
    ((geometric-select alpha (+ start 1)) (- 1 alpha))))

(define receding-distribution
  (stochastic-thunk->distribution
    (lambda () (geometric-select 1/4 0))))
```

Figure 1: An infinite distribution

```
(define (tree-structure alpha)
  (discrete-select
    ('() alpha)
    ((cons (tree-structure alpha) (tree-structure alpha)) (- 1 alpha))))

(stochastic-thunk->distribution (lambda () (tree-structure 1/3)))
```

Figure 4: A recursive distribution over tree structures.

```
(stochastic-thunk->distribution
  (lambda () (make-list (geometric-select 1/4 0) 'a)))
```

Figure 2: A distribution over lists whose elements are references to the symbol a, and whose lengths are distributed according to the geometric distribution.

```
(stochastic-thunk->distribution
  (lambda ()
    (map (lambda (ignore)
          (make-list (geometric-select 1/4 0) 'a))
         (make-list (geometric-select 1/2 0))))))
```

Figure 3: A distribution over lists of lists, all of unbounded lengths.

8. DISCUSSION AND FUTURE WORK

This work has shown how to offer an interface for embedding probabilistic modeling into a full, practical programming language. Probabilistic Scheme is a proof-of-concept implementation demonstrating that this interface is reasonable, and does not require the creation of new programming languages from scratch.

There remain plenty of questions whose answers will help make this system more practically useful:

- How useful is searching possibilities best-first, and what ways are there that the user could supply search heuristics?
- Numerical roundoff error is important since extremely small numbers are known to arise in the practice of probabilistic modeling. What are the right techniques for dealing with it?
- Is it possible to discover no-good sets of discrete selections and avoid them in some dependency directed manner, e.g. [2], [13]?

- What are the right decision-theoretic constructs that naturally force distributions only as far as is useful?
- Can Probabilistic Scheme be extended to continuous probability distributions?
- Can Probabilistic Scheme be extended to reasoning over first-order and other more general propositions, rather than just distributions?

9. ACKNOWLEDGMENTS

The author would like to thank Yu-hsin Chen and Taylor Campbell for assistance with the implementation; and Gregory Marton, Brian Milch, Gerald Sussman and Boris Katz for many useful discussions.

10. REFERENCES

- [1] H. Abelson, G. J. Sussman, and J. Sussman. *Structure and Interpretation of Computer Programs*. MIT Press, 1984.
- [2] F. Bacchus, S. Dalmao, and T. Pitassi. Value elimination: Bayesian inference via backtracking search, 2003.
- [3] N. Friedman, L. Getoor, D. Koller, and A. Pfeffer. Learning probabilistic relational models. *Proceedings of the Sixteenth International Joint Conference on Artificial Intelligence*, pages 1300–1309, 1999.
- [4] R. Kelsey, W. Clinger, J. Rees, et al. Revised5 Report on the Algorithmic Language Scheme. *SIGPLAN Notices*, 33(9):26–76, 1998.
- [5] D. Koller and A. Pfeffer. Probabilistic frame-based systems. *Proc. AAAI*, 98, 1998.
- [6] J. McCarthy. A Basis for a Mathematical Theory of Computation. In P. Braffort and D. Hirschberg, editors, *Computer Programming and Formal Systems*. North-Holland, Amsterdam, 1963.
- [7] B. Milch, B. Marthi, S. Russell, D. Sontag, D. Ong, and A. Kolobov. BLOG: Probabilistic models with unknown objects. *Proceedings of the Nineteenth Joint Conference on Artificial Intelligence*, 2005.

- [8] A. Pfeffer. IBAL: A probabilistic rational programming language. In *IJCAI*, pages 733–740, 2001.
- [9] S. Russell and P. Norvig. *Artificial intelligence: a modern approach*. Prentice-Hall, Inc. Upper Saddle River, NJ, USA, 1995.
- [10] M. Sahami, S. Dumais, D. Heckerman, and E. Horvitz. A Bayesian approach to filtering junk e-mail. *Learning for Text Categorization: Papers from the 1998 Workshop*, 62, 1998.
- [11] E. Segal, B. Taskar, A. Gasch, N. Friedman, and D. Koller. Rich probabilistic models for gene expression. *Bioinformatics*, 17(Suppl 1):S243–52, 2001.
- [12] D. Spiegelhalter, A. Thomas, N. Best, and W. Gilks. BUGS: Bayesian inference using Gibbs sampling. *MRC Biostatistics Unit, Cambridge, England*. www.mrc-bsu.cam.ac.uk/bugs, 1994.
- [13] R. M. Stallman and G. J. Sussman. Forward reasoning and dependency directed backtracking in a system for computer-aided circuit analysis. AI Memo 380, MIT Artificial Intelligence Laboratory, September 1976.
- [14] S. Thrun, M. Montemerlo, and A. Aron. Probabilistic Terrain Analysis For High-Speed Desert Driving. *Proc. Robotics Science and Systems, Philadelphia, PA, USA, August*, pages 16–19, 2006.